

Subsegmental language detection in Celtic language text

Abstract

Bilingual communities contribute majorly to the code-switching phenomena, where speakers produce sentences containing words and expressions from a second language. With the increase in the data on the web, specially on informal platforms like that of the social media, it becomes increasingly important to detect such irregularities and handle them appropriately. In this paper we propose ways of dealing with code-switching and detecting sub-segments for texts in three Celtic languages: Breton, Irish and Welsh.

1 Introduction

Determining the language of a piece of text is one of the first steps that must be taken before proceeding with further computational processing. This task has received a substantial amount of attention in recent years (Cavnar et al., 1994; Lui and Baldwin, 2012). However, previous research has on the whole assumed that a given text will be in a single language. When dealing with text from formal domains, this may be the case — although there are exceptions — such as quotations embedded in the text in another language. But when dealing with informal text, particularly in languages where the speech community is predominantly bi- or multi-lingual, this assumption may not hold.

There are several ways in which an informal text may contain parts in different languages, such as code switching, quotations, named entities, interjections, translations, etc. Some of these are presented in Table 1.

The work presented in this paper was motivated by the problems in normalising non-standard input for the Celtic languages as a precursor to machine translation. When applying a normalisation strategy to a piece of text, it is necessary to first know the language of the piece of text you are applying it to.

The remainder of the paper is laid out as follows. In section 3 we describe the problem in more detail and look at relevant prior work before proposing a novel method of sub-sentential language detection. Section 4 describes the evaluation methodology. Then in section 5 we present the results of our method and compare it against several other possible methods. Finally, section 6 presents future work and conclusions.

Code switching:	You're a [Meiriceánach, cén fáth] are you [foghlaim Gaeilge?!]
Quotations:	The anthem starts with the words ['Mae hen wlad fy nhadau...']
Named entities:	[Dr Jekyll] ha [Mr Hyde] embannet gant [Éditions Aber]
Interjections:	Hey, that's great, [diolch yn fawr!]
Translations:	Bloavezh mat d'an holl ! [Bonne année à tous !]

Table 1: Some examples of text segments containing more than one language

2 Related Work

Code-switching and segment detection problems have been the subject of previous research. A good deal of work has been done on detecting code-switched segments in speech data (Chan et al., 2004; Lyu

et al., 2006). It is seen that language modelling techniques have shown promise earlier, such as in Yu et al. (2013) the experiment on Mandarin-Taiwanese code-switched sentences show a high accuracy in terms of detecting code-switched sentences.

Chan et al. (2004) has made use of the bi-phone probabilities and calculated them to measure a confidence metric, to (Lyu et al., 2006) which has made use of a classification method, named syllable-based duration classification, which uses the tonal syllable properties along with the speech signals to help predict the code switch points. Yeong and Tan (2010) uses syllable structure information to identify words in code-switched text in Malay-English, however they did not recognise segments in running text, only identifying individual words.

3 Methodology

We use the character n-gram approach along with some heuristics which are relevant to our problem domain of identifying segments for subsequent processing. We would like to both predict the code switched points but looking at the surrounding structure also decide the inclusion of them into the current or the next segment.¹

3.1 Alphabet n-gram approach

We first built character language models using IRSTLM (Federico et al., 2008) for the five languages in question. For English and French a model was trained using the EuroParl corpus (Koehn, 2005). For Breton, Welsh and Irish we used corpora of text crawled from the web. For each of the languages we sampled 1.5 million words.

In order to build a character language model as opposed to the standard word-ngram based model we replaced spaces with the underscore symbol ‘_’, and then placed a space character between each character. For example, ‘sláinte’ would be broken down into a sequence of {‘_ s’, ‘s l’, ‘l á’, ‘á i’, ‘i n’, ‘n t’, ‘t e’, ‘e _’}. Then the probability for the alphabet sequence is calculated for each language to be used in the experiment.

3.2 Sequence chunking

This section describes the heuristics taken into account while performing the chunking of the input data according to the sequences of segments in different languages.

A flag is set at the beginning of the input text, and this moves forward to the position in the text upto where confident chunking has been performed. For example, in case of a continuous prediction there would be no change and since the predicted tag of the new word corresponds to that already of the expanding chunk, this word is also included in the same chunk. In case of a doubt, the flag is set at the same position, and no final decision on labelling the word is performed, to have more confidence on the code-switch, the process moves on to the next token, if it corresponds to that of the previous chunk, then both these undecided words are labelled and included in the chunk, otherwise a change is noted and a new chunk is created with the changed label. In this process, we are trying to identify the sub-segments that we need, thus performing both tasks at once, the first being language prediction and second being the chunking decision.

3.3 Word-based prediction

This is a simple heuristic which was designed on the basis of the the most common words in the wordlists of the languages which are in question. After checking each word against both of the word lists,² it is associated to one language or another. In case of a conflict, for example, when the word exists in both wordlists, or in the case that it is unknown to both, the option of continuing with the previous span was taken and the previous selected tag was labelled, thus increasing the chunk.

¹The software used in this experiment is available online at: `removedforreview`.

²For the wordlists we used the *aspell* wordlists widely available on Unix systems.

```
[en You're a] [ga Meiriceánach, cén fáth] [en are you] [ga foghlaim Gaeilge?!]
@steffdafydd [cy gorfod cael bach o tan] [en though init]
[en omg] [cy mar cwn bach yn] [en black and tan] [cy a popeth,] [en even cuter!!]
```

Figure 1: Example of text from a microblogging site chunked

3.4 Word-based prediction with character backoff

A better way to predict the spans of the sub-segments of the text is to include the two methods of word and character-based techniques as described above. In case of the word being present in only one of the two monolingual word lists the classification is simple, but in case of a conflict, a character bigram backoff was introduced to help us disambiguate the language label. This method works well because the earlier heuristic approach of just labelling the word with the label of the span which is expanding, would mean less code-switch detection and more shift towards the majority class.

4 Evaluation

We hand-annotated a small evaluation set from a selection of posts to a popular microblogging site³. The posts ‘tweets’ were filtered into three sets which had been identified as Irish, Welsh and Breton. Certain tokens were escaped from the data, such as mentioned (which are preceded by an @ symbol), subject tags ‘hashtags’ which are preceded by a # symbol), hyperlinks and the sequence `rt` which stands for ‘re-tweet’. An example of the content of our corpus after annotation is given in Figure 1. All of the tweets had at least one instance of code-switching.

For the Evaluation procedure, we follow the footsteps of the CoNLL-2000 shared task on language-independent named entity recognition: dividing text into syntactically related non-overlapping groups of words. This chunking mechanism (Tjong Kim Sang and De Meulder, 2003) is very similar to ours, in terms of words which only belong to one category (here, language), and also evaluation based on the segment structure present in the data. The chunks here are such that they belong to only one language.

For our current task, each of the texts have been limited to two languages i.e. the primary language (the Celtic language) and the secondary language (the ‘majority’ language). Hence the type of chunking tags are limited to these. The evaluation statistics shown in Tables 3 and 4 mention two values for each of the experiment conducted on the three bilingual language datasets. The first, is the percentage of correctly detected phrases, which is the overall precision and the second is the number of phrases in the data that were found by the chunker, which is the overall recall.

Apart from the techniques discussed in section 3, some baselines are also used to give a comparative view of how well all the mechanism perform.

4.1 Baseline

This is the most naïve method of classification, we used the language identification tool `langid.py` (Lui and Baldwin, 2012) on the whole dataset and labelled all the individual lines according to this single majority classification. As no chunking is performed we can expect that the precision and recall will be very low. However it does provide a reasonable baseline for the per-word accuracy.

4.2 `langid` character bigram and trigram prediction

After restricting the predicted languages to be amongst the two in the dataset, we used the character bigram and then trigram probabilities to predict the detected language for each token, some rules like in section 3 were followed which included that the span of a segment or a code-switched phrase is more than a word, the social media non-language entities such as the hashtags, usernames and urls do not make a difference.

Place licence statement here for the camera-ready version, see Section ?? of the instructions for preparing a manuscript.

³<http://indigenoustweets.blogspot.in/2013/12/mapping-celtic-twittersphere.html>

Pair	Language	Statistics (%)	
		Tokens	Segments
Irish—English	Irish	332	40
	English	379	42
Welsh—English	Welsh	419	64
	English	378	66
Breton—French	Breton	388	54
	French	379	53

Table 2: Document statistics of the annotated data used.

System		Irish—English		Welsh—English		Breton—French	
		Irish	English	Welsh	English	Breton	French
baseline	<i>p</i>	2.50	0.0	0.0	0.0	0.0	0.0
	<i>r</i>	2.56	0.0	0.0	0.0	0.0	0.0
langid2	<i>p</i>	0.00	7.14	0.0	18.18	0.0	28.30
	<i>r</i>	0.00	5.00	0.0	14.12	0.0	18.99
langid3	<i>p</i>	5.00	14.29	0.0	21.21	1.85	20.75
	<i>r</i>	5.41	8.45	0.0	14.58	1.92	12.36
wordlist	<i>p</i>	32.50	28.57	26.69	40.91	57.41	33.96
	<i>r</i>	23.64	26.09	26.03	33.75	47.69	33.33
bigram	<i>p</i>	32.50	35.71	23.44	19.70	57.41	52.83
	<i>r</i>	22.41	26.79	15.31	16.67	41.33	37.84
wordlist+bigram	<i>p</i>	52.50	50.00	32.81	31.82	70.37	67.92
	<i>r</i>	38.18	43.75	24.14	25.61	57.58	57.14

Table 3: Precision, *p* and recall, *r* for the systems by language.

5 Results

As described in the 3 the data collected from Twitter for the three language pairs, was processed using the techniques mentioned. The statistics of the same are given in Table 2. //

In terms of per-word accuracy, the baseline performs as expected, given that half of the words in the test set are from a given language. The `langid.py` based methods allow for segmenting the text into chunks, but favour the majority language.⁴ While the precision and recall are low for the remaining models, we see that we are able to improve the performance by combining the wordlist based model with a character bigram model. And what is more, we are able to begin to not only identify particular words in a language, but also segments.

⁴This could be because of the default models being used.

System	Accuracy (%)		
	Irish—English	Welsh—English	Breton—French
baseline	42.76	42.16	44.07
langid2	50.35	42.28	45.11
langid3	57.24	45.92	43.16
wordlist	79.75	74.28	83.96
bigram	81.29	65.62	76.79
wordlist+bigram	85.79	72.40	88.79

Table 4: Accuracy of the systems over the three language pairs. The accuracy measures how often a token was assigned to the right language, independent of span.

6 Conclusions

This paper has presented a very preliminary investigation into subsegment language identification in Celtic language texts. We have proposed a model that chunks input text into segments, and performs language identification on these segments at the same time. Precision and recall are low, leaving a lot of room for further work. Some of our plans are as follows. We would like to annotate more test data, at least 100 tweets in each language. We would also like to attempt our method using higher order character n-gram models for backoff, and n-gram word language models for detection.

Acknowledgements

References

- William B Cavnar, John M Trenkle, et al. 1994. N-gram-based text categorization. *Ann Arbor MI*, 48113(2):161–175.
- Joyce YC Chan, PC Ching, Tan Lee, and Helen M Meng. 2004. Detection of language boundary in code-switching utterances by bi-phone probabilities. In *Chinese Spoken Language Processing, 2004 International Symposium on*, pages 293–296. IEEE.
- M. Federico, N. Bertoldi, and M. Cettolo. 2008. IRSTLM: an open source toolkit for handling large scale language models. In *Proceedings of Interspeech, Brisbane, Australia*, pages 1618–1621.
- Philipp Koehn. 2005. Europarl: A parallel corpus for statistical machine translation. In *MT summit*, volume 5, pages 79–86.
- Marco Lui and Timothy Baldwin. 2012. *langid.py*: An off-the-shelf language identification tool.
- Dau-cheng Lyu, Ren-yuan Lyu, Yuang-chin Chiang, and Chun-nan Hsu. 2006. Language identification by using syllable-based duration classification on code-switching speech. In *Chinese Spoken Language Processing*, pages 475–484. Springer.
- Erik F Tjong Kim Sang and Fien De Meulder. 2003. Introduction to the CoNLL-2003 shared task: Language-independent named entity recognition. In *Proceedings of the seventh conference on Natural language learning at HLT-NAACL 2003-Volume 4*, pages 142–147. Association for Computational Linguistics.
- Yin-Lai Yeong and Tien-Ping Tan. 2010. Language identification of code switching malay-english words using syllable structure information. *Spoken Languages Technologies for Under-Resourced Languages (SLTU’10)*, pages 142–145.
- Liang-Chih Yu, Wei-Cheng He, Wei-Nan Chien, and Yuen-Hsien Tseng. 2013. Identification of code-switched sentences and words using language modeling approaches. *Mathematical Problems in Engineering*, 2013.