

Finite-state morphological transducers for three Kypchak languages

Jonathan North Washington[†], Ilnar Salimzyanov[‡], Francis M. Tyers^{*}

[†]Departments of Linguistics and Central Eurasian Studies
Indiana University
Bloomington, IN 47405 (USA)
jonwashi@indiana.edu

[‡]Institut für Maschinelle Sprachverarbeitung
Universität Stuttgart
Stuttgart (Germany)
ilnar@ilnar

^{*}Departament de Llenguatges i Sistemes Informàtics
Universitat d'Alacant
E-03071 Alacant (Spain)
ftyers@dlsi.ua.es

Abstract

This paper describes the development of free/open-source finite-state morphological transducers for three Turkic languages—Kazakh, Tatar, and Kumyk—representing a language from each of the three subbranches of the Kypchak branch of Turkic. The finite-state toolkit used for the work is the Helsinki Finite-State Toolkit (HFST). This paper describes how the development of a transducer for each subsequent language took less time. An evaluation is presented which shows that the transducers all have production-level coverage—around 90%—on freely available corpora of the languages, and high precision and recall over a manually verified test set.

Keywords: Kazakh, Tatar, Kumyk, morphology, transducer

1. Introduction

This paper describes the development of morphological transducers for three closely related languages: Kazakh, Tatar, and Kumyk.

The transducers for these languages

2. Languages

The three languages for which transducers were developed belong to the Northwestern branch of Turkic, which is often referred to as the Kypchak branch. This branch can be divided into three subbranches. Kumyk is a member of the Western Kypchak group, Tatar is a member of the Northern Kypchak group, and Kazakh is a member of the Southern Kypchak group (Johanson, 2006, 82-83). As such, each of these three languages represents a different one of the three branches of Kypchak. The geographic distribution of the languages is shown in map 1.

These languages have different amounts of linguistic influence from other Turkic branches (e.g., moderate Oghuz (SE) influence in the Western group, slight Oghuz influence in the Northern group) and from Mongolic languages (moderate influence on the Southern group, lighter in the other groups), and all have heavy influence from Persian.

2.1. Kazakh

Kazakh /qazaq/ is spoken primarily in Kazakhstan, where it is the national language, sharing official status

with Russian as an official language. Large communities of native speakers also exist in China, neighbouring Central-Eurasian republics, and Mongolia. Estimates of the total number of speakers range from 8 million (Lewis et al., 2013) to 11 million (Nationalencyklopedin, 2013) people.

2.2. Tatar

Tatar /totar/ is spoken in and around Tatarstan by approximately 5.4 million people (Lewis et al., 2013). It is co-official with Russian in Tatarstan — a republic within Russia. A majority of native speakers of both languages are bilingual in Russian.

2.3. Kumyk

Kumyk /qumuq/ is spoken in Dagestan, a Republic of the Russian Federation, where it is co-official with a number of other languages of Dagestan (Lewis et al., 2013). There are approximately 430 thousand speakers (Lewis et al., 2013).

3. Background

3.1. Morphological transducers, previous work

The objective of a morphological transducer is twofold: firstly to take surface forms (e.g., алдым) and generate all possible lexical forms, and secondly to take lexical forms (e.g., ал<v><tv><ifi><p1><sg>, алд<n><px1sg><nom>, etc.) and generate one or more surface forms. Since they are implemented as finite-state automata, they are reversible by default.



Map 1: The three sub-branches of Kypchak (North, South, West), roughly divided with black lines, showing the geographic distribution of the three languages for which transducers were developed (tat, kaz, kum). Language codes are from ISO 639-3.

The transducers were designed based on the Helsinki Finite State Toolkit (Linden et al., 2011) which is a free/open-source reimplement of the Xerox finite-state toolchain, popular in the field of morphological analysis. It implements both the **lexc** formalism for defining lexicons, and the **twol** and **xfst** formalisms for modeling morphophonological rules. It also supports other finite state transducer formalisms such as **sfst**. This toolkit has been chosen as it – or the equivalent XFST – has been widely used for other Turkic languages, such as Turkish (Çöltekin, 2010), Crimean Tatar (Altintas, 2001), Turkmen (Tantuğ et al., 2006), and Kyrgyz (Washington et al., 2012), and is available under a free/open-source licence.

The authors learned of another Kazakh morphological transducer in existence (Бекманова & Махимов, 2013) only after production-ready version of our transducer was released. We have not yet been able to evaluate this system or compare it to ours.

Creating a morphological transducers in the above-mentioned formalisms simply involves encoding linguistic knowledge about the language in the formalisms. The **lexc** and **twol** formalisms resemble linguistic formalisms, allowing the coders to work with abstractions resembling linguistic categories such lexemes, morphemes, phonemes, and even archiphonemes—as opposed to a raw FST, where input characters are translated to output characters along a graph.

3.2. Description

The transducers are available / under development in apertium’s subversion repository,¹ in the directories apertium-kaz, apertium-tat, and apertium-kum. The revision of the entire subversion repository that the numbers (stem counts, evaluation, etc.) in this paper

represent is r48137.

4. Methodology

4.1. Development effort

The three transducers discussed in this paper are for Kazakh, Tatar, and Kumyk. The Kazakh and Tatar transducers were originally created as part of an experimental Kazakh-Tatar machine translation system in December of 2010. The Kazakh transducer was expanded during Google Code-In 2010 and 2011, and the Tatar transducer was expanded as part of a prototype Tatar and Bashkir machine translation system (Tyers et al., 2012). The Kazakh-Tatar machine translation system, along with the two transducers, was expanded to production-level quality as part of a Google Summer of Code project in 2012 (Salimzyanov et al., 2013).

The Kumyk transducer was developed starting at the beginning of October, 2013 as experiment to see how difficult it would be to extend lessons learned from the development of the Tatar and Kazakh transducers to a related language. While the Kazakh and Tatar transducers took years of part-time work to reach their current production-quality (around 90%) coverage of corpora, the Kumyk transducer only took a couple weeks to reach this level of quality. This paper explores how the development of the Kumyk transducer benefited from knowledge gained from the development of the Tatar and Kazakh transducers.

The morphotactics of Turkic languages are complex enough that even a linguist who is fluent in the language and has a good linguistic understanding of it may not understand how exactly all morphemes combine. Native speakers educated about the morphology of their languages also do not have an explicit knowledge of the complete morphotactics. Hence it often becomes necessary to use fieldwork methodology to elicit the full extent of the morphotactics, be this a linguist with no to little knowledge of a Turkic language

¹<https://svn.code.sf.net/p/apertium/svn/languages/>

working with a native speaker, or a native speaker who understands the extent of what knowledge is necessary to encode in the transducer. When there is no native speaker of a particular language available, the authors have found that information previously encoded about a closely related language or the intuitions of a speaker of a closely related language may be combined with the use of textual corpora to “elicit” information about the morphotactics of a language. Depending on the contents of corpus and chance, this may not result in a completely accurate model, but it is possible to be thorough.

The Kazakh morphotactics were originally developed based on the Kyrgyz transducer, which was co-authored by a linguist who is fluent in and has a good linguistic knowledge of Kyrgyz together with a native speaker of Kyrgyz, and in consultation with another native speaker of Kyrgyz. The initial developer of the Kazakh morphotactic was the same linguist, who is fluent in and has a good linguistic understanding Kazakh. The morphotactics of Tatar were developed for the most part by a native Tatar speaker, who also worked to polish off the morphotactics of the Kazakh transducer.

As the authors found it difficult to locate native speakers of Kumyk, the morphotactics of the Kumyk transducer were developed based on the existing encoded morphotactics of Kazakh and Tatar (and occasionally Kyrgyz), with consultation of corpora, as described above. A dictionary (Бамматов, 1960) and a grammar (Ольмесов, 2000) of Kumyk were also consulted as needed.

4.2. Transducer contents

Each transducer’s lexc source consists of lists of stems, with each stem pointing at a complex continuation lexicon containing the appropriate morphology for the type of stem.

The tagset for each transducer is designed to be compatible with the others. Each transducer consists of about 120 separate tags, of which close to 20 cover the main parts of speech (noun, verb, adjective, adverb, postposition, interjection, etc.). The remaining tags cover morphological subcategorisation for e.g. case, number, person, possession, transitivity, tense-aspect-mood, etc. The tags are represented as multicharacter symbols, between less-than < and greater-than > symbols. The tagset is quite extensive and still not entirely stabilised, so a full listing is not included here. However, the tags are listed in the source code of the transducers, along with comments describing their usage. Table 1 lists the number of stems of the primary categories in each transducer.

Part of speech	Number of stems		
	Kazakh	Tatar	Kumyk
Noun	2463	2692	2588
Verb	1587	1333	262
Adjective	823	852	217
Proper noun	5412	3523	1444
Adverb	197	205	63
Numeral	63	67	45
Conjunction	52	50	15
Postposition	56	46	12
Pronoun	17	17	18
Determiner	42	37	9
Total:	10942	9009	4687

Table 1: Number of stems in each of the categories

4.3. Categorisation and tagset

4.4. Technolinguistic issues

There is a class of phenomena encountered during the development of these transducers, united by the fact that the morphophonology needs information about the stem’s phonological form that is not provided by the orthographical representation. These phenomena were left by Washington et al. (2012) as future work, but have been extensively implemented in the transducers described by the current paper.

Types of these phenomena include acronyms, “irregular” vowel harmony, Russian loanwords, and numerals. A brief overview of each follows, along with examples, followed by details of our solution.

The best way we have found to handle these phenomena in a two-layer FST like those that we use is to add phonological information at the end of stems that need it in the form of special “abstract letters” at the morphotactic stage that trigger phonological processes at the morphophonological stage and are deleted.

Since `two1` performs all operations in parallel, rules triggered by these must make reference to abstract letters that exist on the input tape but are null on the output tape. This complicates existing rules that have been designed to ignore null characters on the output tape. The `two1` rules resulting from the combination of processes simultaneously triggered by and ignoring output-null characters are rather ... uh, *intricate* ... but they seem to work okay? Just don’t poke them too hard.

5. Evaluation

We have evaluated the morphological analysers in two ways. The first was by calculating the naïve cover-

Құдай Өзінің жаратқандарының бәріне қарап, өте жақсы екенін көрді.
Аллаһь Озью яратған затлаһьа қырап, олар бек яхшы экенин гөрген.

Kazakh	Tatar	Kumyk
Құдай<n><nom>	Аллаһь<n><nom>	
Өз<prn><ref><p3><sg><gen>	Озь<prn><ref><px3sp><nom>	
жарат<v><tv><ger_past><pl><px3sp><gen>	ярат<v><tv><gpr_past>	
бәрі<prn><qnt><px3sp><dat>	зат<n><pl><dat>	
қара<v><tv><gna_perf>	қыра<v><tv><gna_perf>	
,<cm>	,<cm>	
—	олар<prn><pers><p3><pl><nom>	
өт<adv>	бек<adv>	
жақсы<adj>	яхшы<adj>	
e<cop><ger_past><px3sp><acc>	э<cop><ger_past><px3sp><acc>	
көп<v><tv><ifi><p3><sg>	гөп<v><tv><past><p3><sg>	
.<sent>	.<sent>	

Table 2: An example of the output of each of the morphological transducers for the same sentence.

age² and mean ambiguity on freely available corpora. The mean ambiguity measure was calculated by performing an evaluation of precision and recall on some smaller, hand-validated test sets.

5.1. Corpora

We tested the coverage of the Kazakh and Tatar analysers over three separate domains: encyclopaedic text,³ news,⁴ and religion.⁵ As there is currently no Wikipedia in Kumyk, we tested only news and religion.⁶

The coverage of each transducer over the various corpora is shown in table 3.

5.2. Precision and recall

Precision and recall are measures of the average accuracy of analyses provided by a morphological transducer. Precision represents the number of the analyses given for a form that are correct. Recall is the percentage of analyses that are deemed correct for a form (by

Language	Corpus	Tokens	Coverage (%)
Kazakh	Wikipedia	18.2M	85.1 ± 1.33
	News	3.2M	-
	Religion	849K	92.7 ± 1.57
	Average	-	-
Tatar	Wikipedia	128K	-
	News	4.6M	-
	Religion	205K	88.7 ± 0.5
	Average	-	-
Kumyk	Wikipedia	-	-
	News	287K	91.4 ± 0.7
	Religion	227K	92.0 ± 1.1
	Average	-	91.7 ± 0.7

Table 3: Corpora used for naïve coverage tests

²Naïve coverage refers to the percentage of surface forms in a given corpora that receive at least one analysis. Forms counted by this measure may have other analyses which are not delivered by the transducer.

³The following Wikipedia dumps were used: kkwiki-20131006-pages-articles.xml.bz2, FIXME.

⁴All content from <http://www.azattyq.org/> for 2010 was used for Kazakh, as well as all content from 2005 to 2011 on <http://tat.tatar-inform.ru> for Tatar.

⁵We used a Kazakh bible translation available from <https://kkitap.net/> and a Tatar translation of the New Testament available from <http://ibt.org.ru>

⁶The bible corpus is from <http://ibt.org.ru/> and the news corpus consists of all Kumyk content from <http://sh-tavisi.etnosmi.ru/>.

comparing against a gold standard) that are provided by the transducer.

To calculate precision and recall, it was necessary to create a hand-verified list of surface forms and their analyses. We extracted 1,000 unique surface forms at random from a news corpus for each language, and checked that they were valid words in the languages and correctly spelt. Where a word was incorrectly spelt or deemed not to be a form used in the language, it was discarded and a new random word selected.

This list of surface forms was then analysed with the most recent version of the analyser, and each analysis was checked. Where an analysis was erroneous, it was removed; where an analysis was missing, it was

Language	Precision	Recall
Kazakh	98.61	57.98
Tatar	95.03	85.65
Kumyk	-	-

Table 4: Precision and recall

added. This process gave us a ‘gold standard’ morphologically analysed word list of 1,000 surface forms with their analyses. The list is publically available for each language.

We then took the same list of surface forms and ran them through the morphological analyser once more. Precision was calculated as the number of analyses which were found in both the output from the morphological analyser and the gold standard, divided by the total number of analyses output by the morphological analyser.

Recall was calculated as the total number of analyses found in both the output from the morphological analyser and the gold standard, divided by the number of analyses found in the morphological analyser plus the number of analyses found in the gold standard but not in the morphological analyser.

The results for precision and recall will be presented in table 4.

6. Future work

One direction for future work is to develop transducers for more languages. We have already constructed usable prototype transducers for three other Kypchak languages: Bashkir (N), Nogay (S), and Karakalpak (S). Since our ability to develop production-ready transducers is limited by availability of resources, including corpora in the languages and native-speaker consultants, the Western Kypchak languages (aside from Kumyk) have been more neglected by our team. However, these language communities would benefit from computational tools for their languages, and work on them may be bootstrapped from the existing transducers, so working on morphological transducers for these languages is also a priority.

7. Conclusions

We have described morphological transducers for three Kypchak languages—one from each branch of Kypchak—including the development process and performance of the analysers.

Acknowledgements

We would like to thank the Google Code-in (2011) for supporting the development of the Kazakh transducer, and in particular the effort by Nathan Maxson. We would also like to thank the Google Summer of Code (2012) for supporting the development of both the Kazakh and the Tatar transducers.

The authors would also like to express their gratitude to Aida Sundetova and X for assistance in evaluating precision and recall.

References

- Altintas, K. (2001). A morphological analyser for Crimean Tatar. *Proceedings of Turkish Artificial Intelligence and Neural Network Conference*.
- Johanson, Lars (2006). History of Turkic. In Lars Johanson & Éva Á. Csató (Eds.), *The Turkic Languages*, New York: Routledge, chap. 5, pp. 81–125.
- Lewis, M. Paul, Simons, Gary F., & Fennig, Charles D. (Eds.) (2013). *Ethnologue: Languages of the World*. Dallas, Texas: SIL International, seventeenth edn. <http://www.ethnologue.com>.
- Linden, Krister, Silfverberg, Miikka, Axelson, Erik, Hardwick, Sam, & Pirinen, Tommi (2011). *HFST—Framework for Compiling and Applying Morphologies*, vol. Vol. 100 of *Communications in Computer and Information Science*, pp. 67–85. ISBN 978-3-642-23137-7.
- Nationalencyklopedin (2013). kazakiska. In *Nationalencyklopedin*. <http://www.ne.se/>.
- Salimzyanov, Ilnar, Washington, Jonathan North, & Tyers, Francis M. (2013). A free/open-source Kazakh-Tatar machine translation system.
- Tantuğ, A.C., Adalı, E., & Oflazer, K. (2006). Computer analysis of Turkmen language morphology. *Advances in natural language processing, proceedings (Lecture notes in artificial intelligence)*, pp. 186–193.
- Tyers, Francis, Washington, Jonathan North, Salimzyan, Ilnar, & Batalov, Rustam (2012). A prototype machine translation system for Tatar and Bashkir based on free/open-source components. In *Proceedings of the First Workshop on Language Resources and Technologies for Turkic Languages at the Eight International Conference on Language Resources and Evaluation (LREC’12)*. Istanbul, Turkey.

Washington, Jonathan North, Ipasov, Mirlan, & Tyers, Francis M. (2012). A finite-state morphological analyser for Kyrgyz.

Çöltekin, Çağrı (2010). A freely available morphological analyzer for Turkish. *Proceedings of the 7th International Conference on Language Resources and Evaluation (LREC2010)*, pp. 820–827.

Бамматов, З. З. (1960). *Русско-кумыкский словарь*. Москва: Государственное издательство иностранных и национальных словарей.

Бекманова, Г. Т. & Махимов, А. (2013). Графематический и морфологический анализатор Казахского языка. pp. 192–200.

Ольмесов, Нурамат Хайруллаевич (2000). *Сопоставительная грамматика кумыкского и русского языков*. Махачкала: ИПЦ ДГУ.